

Performance Evaluation project: Optimizing cars' trajectory with AI

Ottavy Macéo, Longatte Mathieu, Louison Mocq

Part I

Introduction

The goal of this project is split into five parts:

- Creating racing car environment to simulate simple 2D racing car model.
- Implementing Deep Q-learning and Genetic algorithms to optimize the behaviour of a car on tracks so that the car can have the best trajectories possible.
- Evaluate the performances of Deep Q-learning and Genetic algorithms and compare them.
- Evaluate the performances of Deep Q-learning depending of the hyperparameters.
- As a bonus: evaluate the performance of our best car's behaviour.

All the code have been made with python.

Part II

Car Racing environment

1 Tracks

A track is originally a .png file wich look like the left image of figure 1. Then, the image is converted to a matrix T such that $T[0][0]$ is the bottom left corner. After that, we crop the image, compute the starting point and the lines of track (that will be explained in the reward part) to have a final result which look the right image of figure 1.




Figure 1: .png and computed track

2 Cars' physics

The Car physic is really simple. It is a 2D cartoon-like physics that act as follow:

The car has to main informations: its speed $\in [0, \text{MaxSpeed}]$ and its rotation $\in [0, 360]$. The physics is simple, at each times step, the car move to next coordinates on direction of the car's rotation and of a lenght equal to the car's speed.

If the coordinates of the car is (x, y) , its speed is s and its rotation is α , then, after a time step, the coordinate of the car will be:

$$(s \cdot \cos(\frac{\pi}{180}\alpha) + x, s \cdot \sin(\frac{\pi}{180}\alpha) + y)$$

Moreover, at each time step, the car can make some actions:

- it can accelerate, this will increase the car's speed by a constant
- it can brake, this will decrease the car's speed by reduce the car speed by a constant. The car cannot have a negative speed.
- it can turn, i.e. add a constant $\in [-K, K]$ to its rotation. K is a constant that is the maximum angle the car can turn per each time step.

The behaviour of the car will need to interact with the track therefore we need to decide what is the state of a car, i.e. how the car see the environment. We could give to our algorithms the track matrices, the informations of the car but this will lead to too many arguments because a track can have size 900×600 . Therefore we will need to train on all possible states which will be at least $2^{900 \times 600}$. Therefore, we decided to give a more realistic state which represents how a car racer sees. Then the state of a car is an array of size 8.

- T_0 is the current speed of the car
- $\forall i \in \{1, \dots, 7\}$, T_i is the distance of the car to the next wall in the direction $\alpha + A_{i-1}$ where α is the current rotation of the car and $A = [60, 40, 20, 0, -20, -40, -60]$

Then, the representation looks like figure 2.




Figure 2: Car state

3 Technical aspects of the environment

To manipulate our environment, we use the python packages `gymnasium` which provide code convention for those types of environments, i.e. environments where at each time step, you have one action to do. The environment has to have some essential functions: `reset()` that resets the environment to be able to do another simulation, `render()` that renders the current state of our environment and the most important one is `step()` that does one step of time, i.e. given an action, the `step()` function figures out if the car has crashed or not, moves the car to its next position and returns the new state of the car, a reward and if the car has crashed.

Our environment has a variable named `time` which gives us the opportunity to discretise more or less the time.

4 Rewards

For those types of problems where the AI model has to compute a behaviour, the AI model produces something which looks like a function f that takes a car state and returns an action. We need to specify to our AI model when it produces a good action and a bad action, for instance, if a car crashes, we need to punish the AI model.

We do that thanks to a function `reward` implemented in the function `step()` of our environment. The reward is an integer, the bigger it is the better the action was. To punish the car when it does something bad we do:


- If the car crashes, we stop the simulation and return a reward of -500
- If the car is not moving, i.e. has a speed of 0, the reward is -10

For the positive reward, we have automatically computed some track line (represented in right image of figure 1). If the car crosses next line, it has a reward of $+10 \times$ the number of lines it has cross in the good order with this action. If the car cross a line in the wrong order, it means that it has gone backward, therefore, we punished the car with a reward of -200 and we stop the computation.

On top of that, at each time step, we add to the current reward the speed of the car divided by a constant to encourage the car to go fast.

5 Conclusion

After explaining all of this, here is an example in figure 5. We have plot the trajectory of the car, the green color is when the car has accelerate, red color when it brake and yellow otherwise.



list of rounded reward: [10, 1, 2, 12, 13, 13, 13, 14, 24, 14, 13, 12, 13, -497]

Figure 3: Car state

The total reward of a car behaviour is the sum of all reward of a simulation with a car behaviour.

Part III

Deep Q-learning

6 Markovian decision process

We want to describe the interaction between an agent and its environment. In this report, a Markovian decision process is (S, A, T, R) where :

S is a set state in which our agent can be

A is the set of actions our agent can take : all actions $a \in A$ are available from every state $s \in S$

T is the transition function describing what is the next state of our agent after choosing one action :
 $T : S \times A \rightarrow S$

R is the reward function giving $R(s, a, s')$ if we go to s' from s by choosing action a

7 Q value

We can associate to a state a value that reflect the best cumulative reward that we can obtain in the long run if a choose actions optimally starting from its state. We call this value the Q-value of the state : $Q^*(s)$.

8 Deep Q-Learning

Deep Q-Learning is an algorithm that aims at approximating Q^* using a neural network.

Part IV

Genetic algorithms

9 What are genetic algorithms?

Genetic algorithms (GA) are probabilistic algorithms based on natural selection. Therefore, GA takes some populations which are sets of solutions (here a solution is a car's behaviour), select the best solutions thanks to the reward function. Then, it changes the population by adding new random solutions, adding some mutations which are some small variations of a behaviour, adding some cross-over which are the equivalent of natural reproduction. We will repeat this process a fixed number of generations.

10 Markov Chain modelisation

We will now introduce a Markov chain modelisation to genetic algorithm.

We define a markov chaine $(Y_n)_{n \in \mathbb{N}}$ as folowing:

- A state of $(Y_n)_{n \in \mathbb{N}}$ is a population.
- Let y_0 be a special state such that if $Y_n = y_0$ then it means that the population of state Y_n contain an optimal solution.

Now, the sequence of population of genetic algorithm can be discribe with this markov chains. Y_n represent the population at generation n .

Notice that the state y_0 is an absorbing state. In fact if $Y_n = y_0$ then it mean that $P_n[0]$ is optimal. Since we always kept the best solution of previous populations, it means that $\forall n' > n$, we have that $P_{n'}$ contain an optimal solution. Therefore, $\forall n' > n$, $Y_{n'} = y_0$. Moreover, y_0 is the only absorbing state of $(Y_n)_{n \in \mathbb{N}}$.

If we suppose that our mutation and cross-over are made such that a solution x can reach y by a series of a finite number of those operations. Then, all solutions x can reach an optimal value. Then every state y_n can reach state y_0 . The set of all possible state is finite.

Then $\mathbb{P}(Y_n = y_0) \xrightarrow{n \rightarrow +\infty} 1$

Then $\mathbb{P}(\text{The population } P_n \text{ contain an optimal solution}) \xrightarrow{n \rightarrow +\infty} 1$

11 NEAT

Basic genetic algorithms are not efficient enough to compute an optimize behaviour. Therefore, we will use the famous python packages called NEAT. It is an optimized generalized genetic algorithms with dynamic neural network. By dynamic we mean that the algorithm can add or delete some of the nodes of the neural network. The principle of GA stay the same but we have a lot more hyperparameters.

Part V

Performance Evaluation

Once the models are selected (in this case, Deep Q-Learning and Genetic Algorithms), the next step is to compare them using various metrics. For this purpose, we choose the following metric: the average reward after training.

We measure this value across different training durations and varying numbers of tracks used for training the models. To ensure robustness and evaluate potential overfitting, we test the models on tracks that were not included in the training set.

This approach is applied in the context of an AI project where we train AI agents to complete a racing circuit. The goal is for the cars to complete laps as quickly as possible, and the average reward reflects their performance under these conditions.

12 GA VS Q

13 Q hyperparameter

14 Best car